

School of Computing and Information Systems
The University of Melbourne
COMP30027 MACHINE LEARNING (Semester 1, 2019)

Practical exercises: Week 6

Today, we will expect you to be referring to the API¹ for `scikit-learn` (<http://scikit-learn.org/stable/modules/classes.html>) — you should also refer to previous weeks' exercises where necessary.

1. We will use the *Car Evaluation* dataset from the UCI Machine Learning Repository (<https://archive.ics.uci.edu/ml/machine-learning-databases/car/car.data>).

- (a) Load the data into a suitable format for `scikit-learn`, for example:

```
>>> for line in f:
...    atts = line[:-1].split(",")
...     X.append(atts[:-1])
...     y.append(atts[-1])
```

- (b) How many instances are there in this collection? How many attributes, and of what type(s)? What is the class we're trying to predict, and how many values does it take?
 - (c) Are there any missing attribute values? Is there any evidence that this is an artificially-constructed dataset?
 - (d) What happens if we try to build a classifier (using `fit()`) using this data?
2. Unfortunately, `scikit-learn` isn't set up to deal with our attributes in this format.

- (a) Write some functions that transform our **categorical** attributes into **numerical** attributes, by (perhaps arbitrarily) assigning each categorical value to an integer, for example:

```
>>> def convert_class(raw):
...     if raw=="unacc": return 0
...     elif raw=="acc": return 1
...     elif raw=="good": return 2
...     elif raw=="vgood": return 3
```

- (b) Load the dataset again, this time as integers. Observe that we can actually build a model using this data.
 - (c) Split the data into training and test sets².
3. Read up on different implementations of the **Naive Bayes** classifier in `sklearn.naive_bayes`. Which one do you think is most suitable for the dataset we have?

- (a) Train the (default) Naive Bayes model and determine its accuracy on the held-out test set.
- (b) Compare the accuracies of all three different kinds of Naive Bayes classifier (perhaps on a few different train-test splits). Does this accord with your expectations?
- (c) By default, this implementation of Naive Bayes uses **Laplace smoothing**. Turn this off, and see what happens — what is the significance of the reported accuracy?
- (d) What happens if you increase the smoothing parameter instead? Calculate the accuracy for a range of values from 5 to 500. For the very large values, examine the predicted classes for the test instances — what is happening?

¹Note that there are some small differences between the version installed in the labs and the latest stable version.

²You should probably get in the habit of converting to `numpy` arrays, even though `scikit-learn` will often do this for you.

4. The transformation of the data in Q2 implicitly creates ordinal attributes. At first glance, such a strategy does seem reasonable in light of the given values (such as `small`, `med`, `big`). A different strategy would be to replace each categorical attribute having m values with m **binary** attributes, in `scikit-learn` this uses the `OneHotEncoder`³:

```
>>> from sklearn.preprocessing import OneHotEncoder
>>> ohe = OneHotEncoder()
>>> ohe.fit(X)
>>> X_trans = ohe.transform(X).toarray()
```

Note that this transformation should be done before we split the data into training and test sets. (Why?)

- (a) Check the shape of `X_trans` — how many attributes do we have now? Does this correspond to your expectations?
 - (b) Split the dataset comprised of one-hot attributes into train and test sets. Compare the accuracies of the three Naive Bayes models using ordinal attributes with the three models using one-hot attributes: are you surprised? What can we infer?
5. Recall that we built a `DecisionTreeClassifier` in Week 4.
- (a) Do you think the dataset comprised of ordinal attributes, or the dataset comprised of one-hot attributes would be more appropriate for a typical **Decision Tree**? Check the test accuracy of the default Decision Tree model built on each of these datasets.
 - (b) How does the accuracy of the Decision Tree models compare with Naive Bayes on these datasets? Why might this be?
 - (c) (n.b. This step might not be possible in the labs.) The main strategy for visualising a Decision Tree in `scikit-learn` is through the `export_graphviz()` method. Read up on the method, and explore whether the trees created in the previous question are indeed different.
 - (d) Try altering the value of the `max_depth` parameter, between 1 and `None`⁴. Visualise the resulting trees, if you can. Compare the estimated training and test accuracies; is there any evidence that the algorithm is over-fitting (or under-fitting) for trees of certain depths?
6. The **filtering** approach to Feature Selection in `scikit-learn` can be done using `SelectKBest`, for example:

```
from sklearn.feature_selection import SelectKBest, chi2
x2 = SelectKBest(chi2, k=3)
x2.fit(X_train, y_train) # these two statements can be combined into
X_train = x2.transform(X_train) # a single statement via fit_transform()
X_test = x2.transform(X_test)
```

- (a) What happens to the shape of `X_train` and `X_test` now?
- (b) Find out what the best features⁵ were for your dataset, according to χ^2 :

```
for feat_num in x2.get_support(indices=True):
    print(feat_num)
```

- (c) We'll be re-visiting these ideas again — on more interesting data — in later weeks.

³Note that the previous step, of replacing the strings with integers, is a necessary intermediate step for using the `OneHotEncoder`.

⁴It is perhaps not-so-obvious how deep the tree can possibly become, depending on the dataset that we are examining!

⁵Remember, everything is a number, so we need to do some work to get the name back.